

Supplementary Materials for Prior-Guided Bayesian Image Analysis for Physically Calibrated Longitudinal Maize Height

Haoming Wang, Chong Wang, Yawei Li, Cheng-Ting Yeh,
Patrick S. Schnable, and Peng Liu

Iowa State University, Ames, Iowa, USA

Address correspondence to: pliu@iastate.edu

This document provides the detailed implementation specification, causal information-set definition, calibration derivation, uncertainty calibration, candidate-selection audit, annotation provenance, and plant-level endpoint results supporting the longitudinal analysis in the main article. It contains no separate reference list; citations and bibliographic details are given in the main article.

S1 Detailed implementation specification

S1.1 What is and is not Bayesian in the field implementation

Two sequential operations must be distinguished. The field association code maintains a running horizontal position and base row for each plant. It uses a fixed update gain of 0.5 and does not propagate a position covariance, calculate a data-dependent Kalman gain, or retain position particles. This is a fixed-gain state update, algebraically equivalent to a steady-state Kalman-mean recursion after the gain has been specified. It is not a full Bayesian posterior over position.

The height and growth component is Bayesian. A weighted particle approximation represents the joint distribution of latent height and growth rate. Particles are propagated through the transition model, weighted with the inlier/outlier measurement likelihood, normalized, and resampled when the effective sample size crosses its prespecified threshold. Posterior summaries and one-step predictive distributions are calculated from those weighted particles.

The stronger mechanism in which the predicted height distribution re-ranks competing top-keypoint candidates is evaluated in the controlled candidate-ambiguity simulation and in a matched real-image audit. It was not the top-keypoint rule used to produce the primary field-height results. Thus the field implementation supports claims about prior-guided plant association and base-keypoint updating, while the simulation establishes that a height prior can improve top-candidate choice under known ambiguity. The real-image audit tests whether the available natural candidate sets actually contain that ambiguity.

Table S1: Full parameter and implementation specification for the two empirical analysis streams. A dash indicates that the operation was not applied in that stream.

Component	Quantity or operation	2024–2025 daily pole stream	2021 field-height comparison
Observation unit	Plant identity and chronological key	<code>camera_genotype</code> , left-to-right plant track, capture time	<code>rowid</code> plus <code>plantid</code> , field date
Observation unit	Measurement entering height filter	Depth-corrected red-band pole calibration, cm	Archived single-image mask extent <code>height_sam</code> , multiplied by the recorded geometric scale
Reference	In-scene physical definition	Red-band pole; adjacent band edges 1 ft (30.48 cm) apart	Camera-support pole annotation is an installation asset, not the source of <code>height_sam</code>
Reference	Depth factor	$D_{\text{tgt}}/D_{\text{ref}} = 8.5/10.0 = 0.85$	Separate 2021 scale, 0.415886 cm/pixel, zero intercept
Calibration history	Reference selection	Best passing calibration with timestamp no later than the current image; fit RMSE capped at 0.5 in	Not applicable to archived <code>height_sam</code>
Field association	Expected plants	6 per initialized camera view	Archived identities supplied
Field association	Initialization	First complete six-detection frame, sorted left to right	–
Field association	Assignment	Minimum total absolute horizontal distance; maximum normalized distance 0.16	–

Continued on next page

Table S1 continued

Component	Quantity or operation	2024–2025 daily pole stream	2021 field-height comparison
Field association	Position/base update	$m_t = (1 - K)m_{t-1} + Kx_t$, fixed $K = 0.5$; no covariance state	–
Field association	Empirical diagnostic	2018 matches, 48 more than frozen first-frame anchors	–
Particle filter	Latent state	Height h_t and growth rate r_t	Same
Particle filter	Transition order	Propagate h_t using r_{t-1} , then update r_t	Same
Particle filter	Number of particles	5000	5000
Particle filter	Base random seed	20260723; deterministic plant-specific seed from SHA-256 of plant label and base seed	20260724; same deterministic derivation
Particle filter	Initial height	$N(Y_1, \sigma_1^2)$, clipped to $[0, 450]$ cm	Same
Particle filter	Initial growth	$N(3, 2^2)$ cm/day, clipped	$N(0, 8^2)$ cm/day, clipped
Particle filter	Height limits	0–450 cm	0–450 cm
Particle filter	Growth limits	–4–15 cm/day	–4–15 cm/day
Particle filter	Height-process SD	16.866 cm/ $\sqrt{\text{day}}$, estimated from depth-corrected 2024 daily increments	8.78 cm/ $\sqrt{\text{day}}$, robust scale of 2021 image increments
Particle filter	Growth-process SD	0.35 cm/day/ $\sqrt{\text{day}}$	0.35 cm/day/ $\sqrt{\text{day}}$
Particle filter	Growth-rate evolution	Random walk ($\rho = 1$)	Nonparametric rate-versus-height field from 2024–2025 pole-stream increments; field adopted before growth noise is added

Continued on next page

Table S1 continued

Component	Quantity or operation	2024–2025 daily pole stream	2021 field-height comparison
Particle filter	Outlier probability π	0.0512	0.065
Particle filter	Outlier SD	70.62 cm	59.04 cm
Particle filter	Resampling rule	Systematic resampling when $ESS < 0.5M$	Same
Particle filter	Posterior interval	Weighted 0.025 and 0.975 height-particle quantiles after the current image	Same
Particle filter	Predictive interval	Numerical 0.025 and 0.975 quantiles of the complete particlewise Gaussian mixture	Same algorithm; field-height comparison emphasizes the posterior interval
Measurement SD	Inlier scale	Prespecified root-sum-of-squares budget: 8-cm pose term, pole-fit term, 1.5 cm per band outside fitted range, and band-regularity term; the total is multiplied by 0.85 after depth correction	$12[1 + 1.5g_t]$ cm, where g_t is the observed gust normalized over the five evaluated dates; range 12–30 cm, fallback 15 cm
Measurement SD	Interpretation	Pilot uncertainty model; process and measurement variation are not separately identified	Retrospective full-window weather normalization; prospective use requires externally frozen normalization limits
Primary comparison	Empirical sample	Next image-derived measurements after each first observation	All 132 records, 33 plants, 12 rows, and 5 dates

Continued on next page

Table S1 continued

Component	Quantity or operation	2024–2025 daily pole stream	2021 field-height comparison
Primary comparison	Cluster uncertainty	Cameras for calibration summaries	20,000 row-bootstrap replicates; seed 20260908 (endpoint seed 20260909)

S2 Causal online definition

Let I_t be the image acquired at time t , and let $\mathcal{I}_{\leq t} = \{I_1, \dots, I_t\}$ denote the accumulated image pool. With detector weights, thresholds, reference definitions, geometry, and filter settings fixed, the operational update is

$$\mathcal{O}_t = A(I_t, \mathcal{O}_{t-1}; \theta_{\text{fixed}}), \quad \mathcal{O}_t \text{ is measurable with respect to } \mathcal{I}_{\leq t}. \quad (\text{S1})$$

Here $\mathcal{O}_t$ contains the association state, physical image measurement, and height-growth particle distribution returned immediately after I_t is added. An image I_{t+k} , $k > 0$, is absent from Equation S1. When it arrives, it initiates a new update; it does not revise $\mathcal{O}_t$.

The field association step predicts each track by carrying forward its running horizontal position m_{t-1}^x and base row m_{t-1}^b . Current detections are assigned to these carried states by minimum-cost matching. For an accepted observation (x_t, b_t) ,

$$m_t^x = (1 - K)m_{t-1}^x + Kx_t, \quad (\text{S2})$$

$$m_t^b = (1 - K)m_{t-1}^b + Kb_t, \quad K = 0.5. \quad (\text{S3})$$

Equation S3 stores means only. The code does not store a position covariance, so m_t^x and m_t^b are called running position states rather than position posteriors in this supplement.

The physical reference attached to a plant image is also selected as of t : only accepted reference fits whose calibration time is no later than the plant-image capture time are eligible. The selected top and base pixels are converted to centimeters, the depth factor is applied, and the resulting measurement updates the height-growth particle distribution. The cached detection-to-measurement audit recomputed 391 camera-time prefixes and compared 51,035 previously emitted rows, with 0 mismatches.

This prefix audit does not imply that every model-development action occurred prospectively. The 2021 analysis normalizes wind over its five observed dates and transfers a growth field learned from 2024–2025. It is therefore a retrospective cross-year evaluation of a sequential recursion. A prospective deployment must freeze the wind-normalization range, growth field, detector, gates, and all hyperparameters before the first test image.

In the simulation-only candidate extension, candidate j with detector evidence c_{tj} and implied

height y_{tj} is scored before the measurement update as

$$j_t^* = \arg \max_j \{c_{tj} p(y_{tj} \mid \mathcal{I}_{\leq t-1})\}. \quad (\text{S4})$$

This operation lets the height prior affect which top candidate becomes the measurement. It is distinct from applying the particle filter after an independently chosen detection, and it was not used to generate the empirical field-height table.

S3 Manual 2021 support-pole annotations

The source annotation directory contains one CVAT XML file for each of 12 camera rows, named `C-XXX.xml`. An XML can contain multiple dated images, and each image can contain several polylines labeled `pole1` through `pole4`. These labels distinguish visible camera-support poles within an image. The numeric suffix does not specify pole depth or physical length.

For each polyline, the importer reads all vertices, takes the smallest image row as the nominal camera-mount endpoint and the largest image row as ground contact, and records the sum of segment lengths in pixels. The physical definition is the nominal vertical distance from ground contact to the camera mounting or optical-center height: 5 ft (152.4 cm). Annotation status is read from the CVAT attribute named `status`. The imported data contain 199 traces: 172 marked usable and 27 marked unusable.

An earlier export contained two filenames for the same C-039 scene. Following the project decision, `C-039_2021-07-30.JPG` was retained and `C-039_2021-07-30JPG.JPG` was removed. The updated `C-039.xml` also includes `LH54.21-C-039_2021-07-12-10-57_CAM827.JPG`. The curated source therefore contains 61 unique XML-referenced images, and every referenced image was found in the source image directory.

The installation records separately specify four horizontal field-layout intervals of 35 in, totaling 140 in between adjacent rows. These are not graduations on the vertical support poles. They also differ from the 1-ft red-band spacing on the 2024–2025 calibration poles. The 2021 import is consequently an installation and annotation resource; it is not pooled with red-band calibration observations.

The importer writes three canonical products:

- `data/processed/manual_poles_2021/manual_pole_annotations_2021.csv`, one row per retained trace;
- `data/processed/manual_poles_2021/pole_definitions_2021.csv`, the physical definitions and source records; and
- `data/processed/manual_poles_2021/summary.json`, counts, duplicate decision, and image-presence audit.

S4 Physical calibration and depth correction

The 2024–2025 red-band pole supplies positions with known 30.48-cm spacing. Within one reference image, band image rows y_j are centered and scaled as $u_j = (y_j - \bar{y})/s_y$, and the calibration curve is

$$z(y) = \frac{au + b}{cu + 1}. \quad (\text{S5})$$

Only height differences are used. Consequently, an arbitrary additive physical datum along the pole does not affect a top-minus-base extent. Held-out bands evaluate interpolation, while the root-direction diagnostic removes lower bands, fits retained upper bands, and predicts downward. The reverse lower-to-upper calculation is not evidence about root extrapolation.

For the depth correction, place a camera at physical height H . Under the ideal pinhole model with negligible roll and a level shared ground plane, the image row of a point at height z and horizontal distance D can be written, up to the image-row sign convention, as

$$y(z, D) - y_{\text{hor}} = f_{\text{px}} \frac{H - z}{D}. \quad (\text{S6})$$

For two points separated by a vertical extent $h = z_{\text{top}} - z_{\text{base}}$, subtraction removes H and the horizon row:

$$|\Delta y(D)| = \frac{f_{\text{px}}}{D} h. \quad (\text{S7})$$

A scale learned at reference depth D_{ref} would assign a target pixel extent the value

$$h_{\text{pole scale}} = \frac{D_{\text{ref}}}{f_{\text{px}}} |\Delta y(D_{\text{tgt}})| = \frac{D_{\text{ref}}}{D_{\text{tgt}}} h_{\text{target}}. \quad (\text{S8})$$

Therefore,

$$h_{\text{target}} = \kappa h_{\text{pole scale}}, \quad \kappa = \frac{D_{\text{tgt}}}{D_{\text{ref}}}. \quad (\text{S9})$$

For the documented 2024–2025 geometry, $D_{\text{tgt}} = 8.5$ ft and $D_{\text{ref}} = 10.0$ ft, so $\kappa = 0.85$. The measurement standard deviation, being in the same linear units, is multiplied by the same factor. The factor is treated as fixed in the reported analysis; uncertainty in measured distances is not propagated.

Equation S9 is exact only within the model assumptions used in the derivation. Camera tilt or roll, lens distortion, uneven ground, different plant depths within a row, uncertain distances, and within-season camera movement can produce residual error.

S5 Height-growth particle filter and exact predictive mixture

For elapsed time Δ_t in days, the implemented transition order is

$$h_t^{(m)} = \text{clip}_{[0, 450]} \left[h_{t-1}^{(m)} + \Delta_t r_{t-1}^{(m)} + \sigma_h \sqrt{\Delta_t} \epsilon_{h,t}^{(m)} \right], \quad (\text{S10})$$

$$r_t^{(m)} = \text{clip}_{[-4, 15]} \left[T \left(h_t^{(m)}, r_{t-1}^{(m)} \right) + \sigma_r \sqrt{\Delta_t} \epsilon_{r,t}^{(m)} \right], \quad (\text{S11})$$




Figure S1: Idealized calibration-transfer geometry. (a) Side view with the 2024–2025 camera at 5 ft, the target row at $D_{\text{tgt}} = 8.5$ ft, and the red-band reference at $D_{\text{ref}} = 10.0$ ft, giving $\kappa = 0.85$. (b) Image-plane construction: the projective curve maps detected top and base rows onto the pole scale before the fixed depth factor is applied. The 30.48-cm labels refer to the red-band calibration pole, not the separate 2021 support poles.

where the innovations are independent standard normal variables. For the daily pole stream, $T(h, r) = r$. For the 2021 transferred-growth-field configuration, $T(h, r) = \hat{r}(h)$, the nonparametric rate-versus-height field estimated from 2024–2025 image increments. Manual 2021 field heights do not enter $\hat{r}$.

Conditional on height particle $h_t^{(m)}$, the observation law is

$$p(Y_t = y \mid h_t^{(m)}) = (1 - \pi_t) \phi(y; h_t^{(m)}, \sigma_t^2) + \pi_t \phi(y; h_t^{(m)}, \sigma_{\text{out}}^2). \quad (\text{S12})$$

Let w_m^- be normalized particle weights after propagation and before observing Y_t . The exact one-step predictive density and cumulative distribution are

$$p_t^-(y) = \sum_{m=1}^M w_m^- \left[(1 - \pi_t) \phi(y; h_t^{(m)}, \sigma_t^2) + \pi_t \phi(y; h_t^{(m)}, \sigma_{\text{out}}^2) \right], \quad (\text{S13})$$

$$F_t^-(y) = \sum_{m=1}^M w_m^- \left[(1 - \pi_t) \Phi\left(\frac{y - h_t^{(m)}}{\sigma_t}\right) + \pi_t \Phi\left(\frac{y - h_t^{(m)}}{\sigma_{\text{out}}}\right) \right]. \quad (\text{S14})$$

Predictive endpoints at probabilities 0.025, 0.05, 0.10, 0.25, 0.50, 0.75, 0.90, 0.95, and 0.975 solve $F_t^-(q_p) = p$ by numerical root finding. These quantiles provide central 50%, 80%, 90%, and 95% intervals. The predictive probability integral transform is $F_t^-(Y_t)$. The reported predictive standard deviation uses

$$\text{Var}(Y_t \mid \mathcal{I}_{\leq t-1}) = \sum_m w_m^- \left(h_t^{(m)} - \bar{h}_t^- \right)^2 + (1 - \pi_t) \sigma_t^2 + \pi_t \sigma_{\text{out}}^2. \quad (\text{S15})$$

Equations S14 and S15 retain both mixture components. They replace the earlier normal approxima-

tion based only on latent-particle and inlier variances.

After Y_t is observed, the weights are updated as

$$w_m \propto w_m^- p(Y_t | h_t^{(m)}), \quad \sum_m w_m = 1. \quad (\text{S16})$$

The posterior mean and interval for latent height are computed from these updated weights. A posterior interval therefore targets h_t after the current image, whereas the predictive interval targets the next noisy image-derived measurement before it is observed. Coverage for those two targets is not interchangeable.




Figure S2: Selected 2025 daily pole-camera trajectories from the field-executed height-growth particle filter. Red points are depth-corrected image measurements, blue lines are posterior means after the corresponding image, and shading is the weighted 95% latent-height posterior interval. Later observations initiate later updates; they do not retrospectively revise the plotted earlier estimates. The figure illustrates posterior behavior against image-derived measurements and is not validation against manual plant height.

S6 Predictive-uncertainty calibration across levels and years

The 0.85 depth conversion changes every quantity expressed in centimeters. The corrected analysis therefore applies it to the image-derived height and its measurement standard deviation, and estimates the process and outlier scales from already corrected 2024 increments. The 2025 records are not used to estimate these values. Over the complete quality-controlled stream, the unit-consistent predictive mean had mean absolute error 15.13 cm, mean log score -4.440, and weighted interval score 8.08 cm.

Table S2 and Figure S3 show calibration at multiple levels. On 372 held-out 2025 forecasts from 2 cameras, the 95% interval achieved 94.9% coverage with a 97.8-cm mean width. The 50%, 80%, and 90% intervals covered 65.1%, 86.3%, and 92.7%, respectively. This pattern is conservative at the lower levels even though the 95% result is near nominal. The small number of test cameras limits claims about general cross-camera calibration.

Table S2: Central predictive-interval calibration. Parameters for the unit-consistent model were estimated from 2024 only. Each target is the next image-derived measurement before it was observed.

Split	Forecasts	Cameras	Nominal (%)	Coverage (%)	Mean width (cm)
2024 development	1266	6	50	63.7	30.7
			80	87.8	59.4
			90	94.1	78.4
			95	96.8	98.2
2025 held-out test	372	2	50	65.1	30.6
			80	86.3	59.3
			90	92.7	78.2
			95	94.9	97.8




Figure S3: Empirical coverage and interval width for the unit-consistent predictive distribution. Parameters were estimated from 2024 development records; 2025 was held out. The diagonal in the coverage panel marks exact nominal coverage. Points above the diagonal indicate conservative intervals.

S7 Matched real-image candidate audit

The real-image audit uses cached candidates from the released 2021 pose checkpoint. Each temporal method receives identical candidates, confidence scores, association gates, fixed scale (0.415886

133 cm/pixel), and particle-filter settings. Human-mask root positions initialize the first image only.
 134 Subsequent associations use the running state; manual field height is reserved for scoring. The
 135 matched subset contains 63 plant-dates from 17 tracks in 6 cameras.

136 The post-detection and prior-in-candidate methods agreed on every selected candidate (Table S3).
 137 Their errors are consequently identical. This is a useful negative control: under these natural-image
 138 candidates and gates, there was no ambiguity for the height distribution to resolve. It cannot verify
 139 the accuracy gain seen in the controlled ambiguity simulation.

Table S3: Real-image candidate audit. Manual field heights were used only for scoring.

Method	Signed error (cm)	MAE (cm)	RMSE (cm)
Single frame	-9.55	13.38	18.01
Identical candidate plus post-detection filter	-9.27	13.07	17.72
Height prior in candidate scoring plus same filter	-9.27	13.07	17.72

140 S8 All-plant mature-date results

141 Table S4 lists the last available 2021 observation for every plant retained in the primary analysis.
 142 No plant is removed using image-minus-field residuals. The single-frame column is the archived
 143 geometrically scaled mask extent; the Bayesian column is the height-growth particle-filter posterior
 144 mean. The interval is the 2.5th to 97.5th percentile of the latent-height posterior after the endpoint
 145 image. Manual endpoint height is used for scoring and is not supplied to the filter.

Table S4: All 33 plants at their last available 2021 field date. Heights
 and interval endpoints are in centimeters.

Plant	Manual field	Single frame	Bayesian mean	Bayesian 95% interval
C_004-p1	110.0	110.8	109.2	[79.2, 142.4]
C_004-p3	132.0	129.8	149.0	[120.3, 178.9]
C_004-p4	140.0	144.5	138.0	[108.4, 169.2]
C_005-p1	144.0	153.9	128.3	[98.5, 158.7]
C_005-p4	130.0	125.6	136.4	[106.8, 167.3]
C_005-p5	138.0	133.9	136.6	[106.7, 167.7]
C_009-p1	148.0	146.8	151.1	[121.6, 182.2]
C_009-p2	132.0	133.5	140.1	[111.9, 170.8]
C_009-p3	136.0	146.0	154.7	[124.6, 185.4]
C_010-p1	116.0	116.0	123.2	[92.1, 157.1]
C_010-p2	112.0	120.8	127.0	[95.9, 161.0]
C_010-p3	70.0	98.1	80.3	[28.5, 136.8]
C_011-p1	196.0	211.9	204.0	[141.2, 239.5]
C_011-p2	196.0	210.4	205.6	[137.7, 244.5]
C_011-p3	196.0	216.7	198.6	[149.6, 231.2]
C_021-p3	48.0	75.9	45.3	[11.6, 79.1]
C_021-p4	110.0	86.9	117.8	[86.8, 159.0]
C_024-p1	116.0	137.2	122.6	[91.1, 157.4]
C_024-p2	112.0	122.7	130.2	[99.5, 164.1]
C_024-p3	118.0	123.9	125.0	[94.8, 158.8]
C_029-p1	210.0	212.1	208.7	[178.5, 238.0]
C_029-p4	200.0	217.9	216.9	[185.9, 246.8]
C_029-p5	190.0	216.1	212.8	[182.7, 242.6]
C_032-p1	130.0	124.6	125.2	[93.8, 158.6]
C_032-p2	120.0	136.4	138.4	[107.2, 170.6]
C_032-p4	140.0	142.6	129.7	[99.2, 164.6]
C_037-p2	90.0	91.9	113.4	[83.5, 145.5]
C_037-p3	100.0	98.1	117.4	[88.8, 148.8]

Continued on next page

Table S4 continued

Plant	Manual field	Single frame	Bayesian mean	Bayesian 95% interval
C_037-p4	112.0	115.8	110.6	[80.6, 144.4]
C_038-p1	142.0	84.8	83.2	[51.9, 116.8]
C_039-p1	120.0	154.7	147.1	[117.6, 176.6]
C_039-p2	110.0	155.1	144.4	[113.8, 176.6]
C_039-p3	140.0	147.2	136.7	[107.0, 167.5]

Across these 33 endpoints, single-frame mean absolute error was 14.45 cm and the Bayesian mean absolute error was 12.38 cm. The Bayesian signed error was +6.17 cm. The row-cluster bootstrap interval for the 2.08-cm paired mean-absolute-error reduction was -1.47 to 6.82 cm. These endpoint results are secondary because the interval is wide and includes zero; the main article uses all 132 plant-date records as its primary comparison.

S9 Data provenance and reproducibility

All paths below are relative to the public GitHub repository.

Table S5: Data and result provenance for the revised submission.

Asset	Canonical path or source	Role and limitation
2021 pole XML	data/manual_annotations/poles_2021/C-*.xml	Copied from authenticated Reports/pole_calibration ; one XML per camera row
2021 pole images	data/raw/pole_calibration_images/	All 61 curated images; confirmed duplicate omitted
Imported 2021 pole table	data/processed/manual_poles_2021/manual_pole_annotations_2021.csv	Trace endpoints, status, pixel length, and physical definition after duplicate removal
Pole definitions	data/processed/manual_poles_2021/pole_definitions_2021.csv	Documents 5-ft support length and distinguishes horizontal 35-in marks from red-band spacing
Daily image manifest	data/processed/pole_camera_daily_manifest.csv	As-of capture records, local/remote identifiers, and checksums
Released pose checkpoint	models/maize_pose_2021_best.pt	2021 YOLO pose weights with SHA-256 checksum and model card
Curated-image pose candidates	outputs/pose_candidates_annotation_set_2021/pose_candidates.csv	Complete cache for all 61 released images
Automatic calibrations	data/processed/automatic_pole_calibrations_daily.csv	Candidate red-band fits and quality-control fields
Causal audit	outputs/online_study/causality_audit_daily.json	Prefix recomputation for cached detections through measurement assembly
Daily Bayesian output	outputs/online_study/pole_camera_bayesian_daily_revised/online_bayesian_height_posteriors.csv	Unit-consistent exact-mixture predictive and posterior summaries

Continued on next page

Table S5 continued

Asset	Canonical path or source	Role and limitation
Daily model settings	outputs/online_study/pole_camera_bayesian_daily_revised/model_hyperparameters.json	Complete saved particle-filter settings
2021 raw comparison	data/raw/heights_compare_2021.csv	Archived image extents and manual field measurements; raw-image-to-extent provenance remains incomplete
2021 canonical posterior	outputs/filter_height_sam_2021/filtered_height_sam__phi1.0_plus_pole_growth_field.csv	Input to all-plants and per-plant summaries
Primary field summary	outputs/filter_height_sam_2021/all_plants_primary.json	All-record and endpoint metrics with row-cluster intervals
Detector-unit correction	outputs/pixel_validation_2021/summary.json	Recomputed centimeter metrics using 0.41588595 cm/pixel rather than the withdrawn field-regression conversion
Field baselines	outputs/field_baselines_2021/	Fixed causal comparators and row-cluster intervals
Real-image candidate audit	outputs/real_image_candidate_ablation_2021/	Matched candidate comparison with human truth used only for scoring
Two-reader packet	annotation/	Blinded sheets, protocol, image manifest, and scoring script; sheets remain unfilled
Generated manuscript values	manuscript/plant_phenomics_numbers.tex and manuscript/plant_phenomics_per_plant_table.tex	Programmatic source for submission numbers.tex and per_plant_table.tex

153 The main regeneration sequence is shown below. The daily-chain shell script requires a Bash environment.
 154 The 2021 filter also requires the locally documented weather workbook used to create its five date-specific
 155 gust values.

```

156 python analysis/import_manual_poles_2021.py \
157 --xml-dir data/manual_annotations/poles_2021 \
158 --image-dir data/raw/pole_calibration_images \
159 --output-dir data/processed/manual_poles_2021
160 bash analysis/run_daily_online_study.sh
161 python analysis/filter_on_height_sam_2021.py
162 python analysis/all_plants_primary_2021.py
163 python analysis/field_baseline_comparison_2021.py
164 python analysis/real_image_candidate_ablation_2021.py
165 python analysis/summarize_uncertainty_revision.py
166 python analysis/recalculate_pixel_validation_units.py
167 python manuscript/generate_plant_phenomics_numbers.py
168 # Regenerate the manuscript figures from the project root:
169 python analysis/plot_online_bayesian_workflow.py
170 python analysis/plot_pole_annotation_2021.py
171 python analysis/plot_calibration_validation.py
172 python analysis/plot_prior_in_detection_sim.py
173 python analysis/plot_field_validation_2021.py

```

174 `python analysis/plot_geometry_diagram.py`

175 The particle-filter output stores the base seed and all model settings. Plant-specific seeds are derived
176 deterministically from the base seed and plant identifier. The all-plants bootstrap uses 20,000 replicates with
177 its seed recorded in `all_plants_primary.json`. The repository includes code and data licenses, a SHA-256
178 manifest, the final available pose checkpoint, and all 61 curated 2021 images. The larger 2024–2025 raw-image
179 archive and the original cross-fitting image set are not distributed; the public release includes the derived
180 measurements needed to regenerate reported summaries. The two independent plant-annotation sheets are
181 prospective assets and must not be described as completed agreement data until both readers finish them.